Eight Tuesdays

Complete process dossier: premise → harness → 96-hour drafting run → human annotation rounds → final book.
Evidence compiled 2026-08-09 from the `autofiction-harness` repo, its embedded book repo, 30+ Codex session rollouts, 3 Claude Code session transcripts, the Claude call ledger, launcher logs, and Apple Books annotation exports. Every number below has a cited source path.

45,090 final words, 24 chapters (premise band was 75,000–88,000)	**62,328** words at fully-autonomous "final assembly" (Jul 5), before any human-directed editing	**–28%** net reduction from machine-final to human-final text
115 + 17 git commits: book repo + harness repo	**406** completed Claude calls (50.6 h of wall-clock inference)	**4,339,906** markdown words generated by the run: a 69:1 scaffolding-to-book ratio
1,179 audit report files (748 Codex-authored, 431 Claude-authored)	**739** chapter revision snapshots in the diffable manifest	**171** human annotations across 3 reading rounds (37 + 127 + 7)

The five text-shaping forces, color-coded throughout

HUMAN direct human input: premise, reading guide, run directives, annotations, taste rules.
RUBRIC the check system: 208-pattern mechanical gate + 79 audit templates + protocols. Built by humans once, executed by models thousands of times.
MODEL autonomous model judgment: drafting, cross-model critique, phase-8 self-revision, cold reads.
GAP evidence that is missing, deleted, or lives on another machine.

Corrections to working memory

- Sessions begin **June 10**, not June 19. Nothing happened on June 19; the run windows are Jun 10–12, Jul 3–5, and editing days Jul 10, Jul 24, Jul 30.
- The harness's formal **editor stage never ran**. `editor/input/` contains only READMEs; there is no `output_edit/`. All editing happened in-place in `drafter/output/` through ad-hoc interactive sessions.
- The book was drafted by **two different Claude models**: Fable 5 (chs 1–22) and Opus 4.8 max (chs 23–24 + all of phase 8 + final assembly), after Fable hit a hard 7-day limit.

Part 0 — Recommended blog flow (everything below, reorganized into narrative order)

Sixteen beats. Each maps to the evidence part of this dossier and to a chart from the statistics appendix (Part 9). Word-budget guidance assumes a long-form post; compress by dropping beats 3, 13, or 14 first.

#	Beat	What it says	Evidence	Chart / pull-quote
1	**Cold open**	Scene, not thesis: Jul 5, 13:13 UTC, the machine writes `final/novel.md`, 62,328 words, judge scores 7–8, "COVERAGE-PRESENT." Five days later you cut 28% of it on taste.	Pt 5.3, Pt 6.3	Stat pair: 62,328 → 45,090; or the funnel (4.34M generated words → 45K book)
2	**The question + platform**	AutoFiction.ai, UMD, "can AI produce a novel worth reading"; the official pipeline's "No human edited the novel text itself" as the foil for your whole approach.	Pt 1.1	Quote the UMD README verbatim
3	**Lineage**	This is book two. VERVE → your annotations on it → verve-legacy audits + the slop casebook. Taste compounds across books.	Pt 1.3, Pt 6.4	Lineage loop diagram
4	**Building the rulebook in a day**	June 9: 17 sources + Wikipedia distilled into 208 patterns / 79 audits / 10 protocols / 15 scaffolds; the Accumulation Principle and safety rails; six of seventeen harness commits are bug fixes to the enforcement machinery itself.	Pt 2	Commit table; tier donut (101 BANNED / 8 CAP / 99 WATCH)
5	**The premise is the real prompt**	~4,300 words of spec carrying all book-specific taste; the mid-run amendment (§9) proving the standard stayed human-owned while the run was live.	Pt 3	Quote §9 verbatim; "That's three words" repair anecdote
6	**Run diary I (Jun 10–12)**	Planning phases in 3 hours; a chapter cell every 1–2 h; doctor rounds; your entire presence = "check" ×20 plus one request that built the snapshot system.	Pt 4	"check" pie; chapter-cell diagram
7	**Hitting the walls**	Five-hour 429s → three-week gap → resume → "cut that to like 22-24 chapters" → the 7-day Fable death mid-chapter-22 → "ok can you use opus 4.8 max instead of fable."	Pt 5.1–5.2	Rate-limit timeline strip; verbatim user quotes
8	**The machine edits itself**	Phase 8: real self-caught defects (class-size canon, "Odile too rewarded"), five batches, three cold-read loops, numeric reconciliation, then "final."	Pt 5.3	Phase-8 table; word counts 61,000→62,328
9	**Reading it like a human**	The Apple Books plumbing (sqlite, epubcfi, asset migration) and round 1's 37 comments: what a reader flags that 79 audits didn't.	Pt 6.0–6.1	Annotation-pipeline diagram; comment-taxonomy table
10	**Four passes, four fingerprints**	Jul 10 roleplay (127 items, taste extrapolated to unread chapters), Jul 10 deslop (–19.5%, "polished but tiring"), Jul 24 casebook (only pass that ADDED words: warmth), Jul 30 deletion-only (7→66 cuts).	Pt 6.2–6.5	Word-trajectory line with authorship bands; churn heatmap
11	**The ledger**	Human vs rubric vs model: what each layer contributed and provably missed; survival-per-chapter; escalation (ch 20: 50 calls, 14+9 doctor rounds, 124 snapshots).	Pt 7	Master per-chapter table (Pt 9.1); survival slope chart

12	**How my prompting changed**	Zero-shot legislation → one-word supervision → plain zero-shot directives → example-first: 1-shot inline, 37/127 labeled annotations, a ~120-example many-shot casebook, then 7-shot + "extrapolate." You independently converged on the harness's own imported principle: demonstration beats specification.	Pt 8 (new)	Examples-per-directive timeline; three-registers table
13	**Model notes**	Why GPT-5.5 orchestrates and Claude writes; the tic war (18 Codex tics in one counterdraft); Fable vs Opus splits; xhigh→high.	Pt 2.4, 5.2	Model-split table
14	**Ecosystem**	UMD pipeline, the bare-harness control experiment you handed off, Fabula at the human-in-the-loop pole.	Pt 11	Comparison table
15	**Honest negatives + gaps**	Length contract failed (45K vs 75–88K); editor stage never ran; self-grading inflation; genre-packaging gap; what you still don't know (reception, bare-harness results).	Pt 7.4, 12.4	"The prior report's claim… was inaccurate" quote
16	**Close**	The thesis the evidence supports: rubrics set the floor (mechanical tells, canon, register), autonomy fills the volume, but taste never automated; it compressed into 171 margin notes and a casebook, and that was enough to move 17,000 words.	Pt 7.1–7.3	Human-input-vs-machine-output ratio stat

Part 1 — The project, the platform, and the two-harness experiment

1.1 Platform context

AutoFiction.ai is a UMD College Park research platform (Chau Minh Pham, Yapei Chang, Mohit Iyyer; paper in preparation) hosting AI-generated novels for community reading, rating, and inline flagging. The official pipeline repo lives locally at `~/Documents/GitHub/autofiction`. Its README makes claims that are the sharpest possible contrast with this project and should anchor the blog post's framing:

> "Human involvement was limited to writing premises, iterating on prompts, debugging pipeline failures, and occasionally restarting failed runs. **No human edited the novel text itself.**" … "A typical 18-chapter novel with 2 review-revision cycles requires approximately 200 agent jobs and costs roughly $120-250 in API usage … All novels published on autofiction.ai were generated using subscription access (ChatGPT Pro and Claude Max, $200/month each)."

Eight Tuesdays inverts that: the last ~17,000 words of net change were human-directed, and roughly 28% of the machine-final text was cut on human taste.

1.2 Detailed harness vs. bare harness

This repo is the maximalist bet: 208 mechanical patterns, 79 audit templates, 10 protocols, 15 scaffolds, coverage matrix, caveats file. The bare-harness counterpart was produced on **Aug 4** by stripping this repo (Codex session `rollout-2026-08-04T18-08-03…`, cwd repo root):

> "ok make a copy of this folder and then with that copy strip out all of the stuff so its a minimal harness. like strip out the checks, any generated artifacts, etc. i basically just want the basic framework of calls and what not so that i can give it to another person and they can build their own harness"

GAP The stripped copy's location and the collaborator's resulting run are not in this repo; document them from your side for the comparison section.

1.3 Lineage: the VERVE book came first

The harness was not built from theory alone. The repo's first commit is `49d9e30 import verve pipeline scaffolding`, and the audit manifest credits "verve-legacy" as the source of a dozen audits (Dialogue Doctor, Comedy Doctor, Positive Prose, Surprise, Anti-Default…). VERVE is a previous novel (an AI-with-700M-users corporate story, 40+ chapters, characters Denny/Karen/Niles/Colton/Dale/Jenna) whose drafting produced the legacy prompts. Crucially, the author's annotations on *that* book produced `author_flagged_slop_examples.md`, the 12.5K-character casebook re-imported on Jul 24 to edit *this* book (Section 6.4). Round-1 annotations also benchmark against it ("great figurative language that I didn't see from the other book"). The lineage is: **VERVE book → author annotations → (a) verve-legacy audit templates baked into this harness, (b) the slop casebook applied manually to Eight Tuesdays.**

Part 2 — Building the harness (one day: June 9)

2.1 The construction commits

The entire check system was assembled in 15 commits on 2026-06-09 (root repo, `git log --reverse`):

#	Commit	What it added
1	`49d9e30`	Import verve pipeline scaffolding + Wikipedia signs source
2	`35f71c8`	Swap Claude model to claude-fable-5 at xhigh effort everywhere
3	`89081f2`	Mechanical pattern registry: 213 rows across 5 TSVs (collation + wikipedia + legacy)
4	`c73f675`	Tiered `quality_gate.py` engine + stage wrappers; fix negation/placeholder/em-dash rows
5	`5bf4db7`	79 judgment audit templates + audit manifest (one audit per source item, owner-split, risk-tiered)
6	`d70a9b8`	10 source-faithful protocol docs + 15 scaffold templates
7	`cf87bc7`	Wire everything into the drafter task prompt
8	`80c2758`	Rewrite editor task prompt, fully genericized
9	`ecf3ac2`	COVERAGE_MATRIX (319+27 rows), CAVEATS.md, README rewrites
10–15	`43660bb…99660df`	Same-day adversarial debugging of the harness itself: gate engine scanning bugs, dead/duplicate/misfiring regex rows, manifest trigger resolution, wrapper lock deadlock/race, launcher limit misclassification, dialogue-manifest segmentation
16	`7f60671` (Jul 10)	Claude effort default xhigh → high (backport: the ledger shows high in live use from Jun 11)
17	`6f72b85` (Jul 10)	Chapter snapshot system: script + launcher watcher + task-prompt discipline; gate reads premise patterns; codex PATH fallback

Story beat: six of seventeen harness commits are bug fixes to the enforcement machinery itself, and the prompt history shows where they came from: the same evening, the human ran `/ultrareview` and `/code-review ultra` on the fresh repo (Jun 9, 18:12–18:41) and then ordered "ok fix all of the bugs, merge, and push everything on main. if u have questions, ask them to me" (19:25). The harness got a multi-agent code review before it ever generated a word. The build prompt itself is also preserved verbatim: "Your task is to set up a new copy of this project in Documents/Github/ named autofiction-harness... I would like to augment/replace those checks/critics with /Users/ndesai-air/Downloads/unified_source_collation.md. This is a complete list of every single thing I want to check. Make sure to add every single item/reccomendation and let me know if you de-duplicate or cannot add it. Do not summarize, be verbose on creating these checks." followed by "Ask me many questions as needed to spec this out and take your time. be thorough." That one instruction is the origin of the COVERAGE_MATRIX's item-by-item traceability and the cannot-add report.

2.2 Where the checks came from

Two sources, traced item-by-item in `COVERAGE_MATRIX.md`:

1. **`unified_source_collation.md`**: a 319-entry collation of **17 sources** (not vendored in the repo; only referenced by line number). Named source tags visible in the matrix and TSV `source_ref` columns: the *BANNED* guide, the *AI-isms Bible*, ten Reddit threads (1: prose dials; 2: em dash; 3: character snapshots; 4: character voice; 5: pacing; 6: challenge/consequences; 7: living world; 8: demonstration beats specification;

9: AI writing workflow; 10: haunted prose field guide), the *LLM writing outputs report*, the *LLM creativity report*, an *Academic novel manual*, a *Novel process report*, and a *Writing ruleset*.

2. **Wikipedia WP:AIWTW** ("Signs of AI writing"), vendored at `checks/sources/wikipedia_signs_of_ai_writing.md` (578 lines).

The matrix uses a disposition vocabulary (IMPLEMENTED-MECHANICAL / -AUDIT / -PROTOCOL / -SCAFFOLD / -PROMPT / MERGED-WITH / CANNOT-ADD / CAVEAT) and includes an explicit **cannot-add report**: publication/submission business, beta readers and critique groups (their anti-sycophancy rules *were* kept as the "Critique stance" header in every audit), writing psychology and daily routines, software recommendations, authorship adjudication, and two Wikipedia-only signs. For the blog: this is the answer to "did you just paste Reddit into a prompt?" No: every item was classified, deduplicated, implemented at a specific layer, or consciously excluded with a reason.

2.3 The four enforcement layers

Layer	Mechanism	Key numbers
Mechanical gate	`checks/quality_gate.py` (581 lines) + 5 TSVs. Curly-quote normalization, code-fence exclusion, scene segmentation, per-line dialogue/narration splitting (incl. unterminated multi-paragraph speech), scopes (narration/dialogue/opening/ending/heading/document), 9 structural engines (one-sentence gravity paragraphs, repeated descriptors, bold density, mixed quotes…)	208 loaded rows: BANNED (blocks commit), CAP (per-scene/chapter/1k-words/book thresholds), WATCH (counted, never blocks; feeds `tics.md`)
Judgment audits	`checks/audits/*.md` + `audit_manifest.tsv`: id, owner (codex inline vs claude via wrapper), tier (core/risk/book), trigger (flags, watch-counts, cadences, calibration gates), output path. Mandatory skeptical-editor stance; a pass without per-criterion quoted evidence, including a nearest-miss quote, is a failed audit	79 templates; core = every chapter; coverage rule: every chapter × every applicable audit ≥ once, verified from filenames
Protocols	`checks/protocols/`: pre/mid/post-draft, required checks, diagnostic questions, revision discipline, editing stages (macro before line), final test, precedence + safety rails	10 docs. Safety rails: no tier may ban em dashes/semicolons/"however" outright; no fix may fake typos or stage messiness; violations are themselves OPEN findings
Scaffolds	`checks/scaffolds/`: master spec, character snapshot (with dialogue samples, word ownership, stress response, "what they get wrong"), world bible, NPC goals, timeline, event-ripple tracker, thread ledger, rolling summaries, canon sheet, continuity log…	15 templates; a stale tracker is a blocking finding at final assembly

Philosophy documents worth quoting in the post: `CAVEATS.md` preserves the sources' own warnings verbatim, above all the **Accumulation Principle** ("No single banned element makes prose bad… The problem is accumulation… density of defaults") which is the entire justification for the three-tier design; and the scope disclaimer that the harness "is not a detector" and must never be used to adjudicate authorship. `checks/protocols/precedence.md` sets the conflict hierarchy: premise/reading guide > stage task prompt > protocols/audits > BANNED/CAP > WATCH.

2.4 Model roles and the cross-model tic war

`drafter/long_novel_task.md` (406 lines) fixes the roles: **Codex (GPT-5.5, xhigh reasoning) is orchestrator/editor-in-chief and may never write chapter prose** ("Codex must not skip Claude's draft and write the chapter alone… do not let Codex backfill prose", even during outages). **Claude drafts every committed chapter.** The stated rationale: "Codex tends toward arrangement, taxonomies, three-item catalogs, and thesis closers; Claude tends toward physical specificity, concrete blocking, and fewer maxims." Tic-policing is symmetric: Codex runs a 13-item Claude-tic audit on every draft (em-dash crutch, shimmer prose, mood verbs "mattered/earned/weighed/carried", "a kind of / something like", recursive negation, profundity-stamp closers, philosopher-narrator intrusions, "lovely/quiet/small" default register…), and Claude audits any Codex counterdraft. The one dual-draft experiment (ch 17, and again ch 22) is decisive evidence for the post: Claude's audit of Codex's ch-22 counterdraft returned **CODEX-TICS-FOUND (18)** ("arrangement-coded object displays, tricolon/status-report dialogue, thesis closers") and the counterdraft was demoted to "pressure only."

2.5 Infrastructure discipline

- **Wrapper-only Claude access** (`scripts/claude_logged_call.sh`, 363 lines): logs every call to `output/logs/claude_calls.tsv`, saves the exact prompt, preserves raw stream JSONL, extracts markdown, serializes calls. Direct `claude -p` is forbidden so provenance is mechanical. 307 prompt files survive under `output/model_prompts/`.
- **Call patience rules** in the task prompt: Fable calls run 12–18 min routinely, chapter drafts 30–60 min; kill only if stream stale ≥5–10 min AND wall clock exceeded. Longest observed call: 59.8 min (ch 4 draft).
- **Launcher** (`drafter/launch.sh`, 626 lines): caffeinated, 96-hour deadline stamped in `.deadline_epoch`, restart classification, chapter-snapshot watcher. 19 Codex attempts logged (38 log files) across the run.
- **Snapshot system**: born from a live user request on Jun 12 (Section 4.6), hardened into the harness Jul 10. 739 manifest rows with labels like `before_dialogue_doctor_round2_repairs`, giving a diffable record of what each audit round did to the prose.

Part 3 — The premise layer HUMAN

3.1 Where the premise actually came from (corrected: it was co-created, and the receipts exist)

The Claude Code transcript for June 9 was pruned by session cleanup, but `~/.claude/history.jsonl` (1,504 prompt records) preserves the full sequence. The premise was **not typed by hand**; it grew from a 56-word seed in about 12 minutes of conversation, in the harness repo, the evening the harness was finished:

> [Jun 9, 20:11] "Describe the steps to create a premise using this repository."
> [20:17] "ok wait describe how to use this repository. i dont have a premise yet. i want to basically cocreate one and refine it with you the ai model. but im confused at the steps to do this."
> [20:39] **"ok cozy low stake romance sort of slice of life set in nyc. something to do with two young people man and women meeting in art gallery and getting mistaken for dating and roped into an art class where feeling blossom. help me develop this idea"**
> [20:49] "your call. write the premise/related stuff but dont launch the entire thing yet. let me see what you come up with"
> [20:57] "was this created by you or another session like Fable 5 xhigh?"

`premise.md`'s file mtime is Jun 9 20:51, between the last two messages. Every load-bearing invention in Part 3's table (Odile knowing from night one, the nested deception, the assignment cards, the diner confession, the 75–88K band, the motif caps) was drafted by the model inside this exchange and ratified by the human the next morning ("ok it looks good, what are the commands to start this process?", Jun 10 12:06). For the blog this reframes the "human premise layer": the human contribution at this stage was a two-sentence idea, a co-creation stance, and approval authority; authorship of the 12.1 KB standard was delegated to the same model that would later be edited against it.

3.1b The premise package as it stands

File	Date	Content
`premise.md` (12.1 KB)	Jun 9 20:51 (evening of the day the harness was finished)	Full story bible: logline, genre contract ("This is a warm book… No villain, no cruelty"), three principals with psychology, the nested two-layer deception + hidden truth (Odile knew all along; must be true throughout, never stated before the showcase), binding story shape (setup, three turns, low-key diner confession, showcase as coda), ~80K length band

`READING_GUIDE.md` (11.1 KB)	Jun 10 22:06 (amended after drafting began)	Form binding (dual POV alternating, June opens and closes, assignment-card epigraphs in Odile's hand); voice differentiation spec ("A reader should identify the POV within a paragraph with the names removed"); comp calibration (Beth O'Leary warmth, "Emily Henry at about 60% banter density", Morisaki Bookshop quiet); plain-speech floor; **zero-tolerance "no dialogue diagnosis" rule**; closed-door heat; motif cap table (hands, light, Winter Kitchen ≤5 uses, first snow falls once); per-character care moves; genre tic catalog; 8 architecture questions the agent must settle; 6 judgment-only quality gates (blind-POV test, reread test, smallness test, warmth floor, earned-stay test, followability test)
`patterns/reading_guide_bans.tsv`	Jun 10 22:04	One mechanically-enforced book ban: `\bload[ –]?bearing\b`, note "Book-specific ban (author order, 2026-06-10)"
`premise_combined.txt` (25.2 KB)	Jun 23 (mid-gap)	Consolidated verbatim package of the three documents

3.2 Mid-run amendment: the human recalibrated the standard while the run was live

The prompt history captures the exact moment the amendment was born. On the evening of day 1 the human was interrogating audit coverage ("how many checks are being done after each chapter", "how many did chapter 2 generate?", "judgements triggered by the dialogue?"), then diagnosed the failure mode personally:

> [Jun 10, 21:13] "imo it's slipping into a common pattern where the characters talk too cleverly for the reader to actually understand what they are trying to convey"
> [21:17] "i mean they shouldn't be diagnosing language like that at all. + add load-bearing or loadbearing to banned words list. and update reading guide. ive stopped it."

That order became READING_GUIDE.md §9, "Mid-run amendments (binding retrofit orders)", and the ban TSV timestamped 22:04–22:06 the same night:

> "**2026-06-10:** The plain-speech floor, the zero-tolerance dialogue-diagnosis rule, and the 'load-bearing' ban were added after chapters 01–03 were committed. Those chapters must be brought into compliance at the next revision touch… Known instances (both confirmed as open BANNED hits by the gate): chapter 01 line ~111 (narration) and chapter 02 line ~87 (June's 'That's load-bearing language for an evening class')…"

Worklog corroboration: Jun 11, ch 4 cycle: "Dialogue Doctor round 3 caught one remaining **amended-reading-guide** blocker: Mara's *'That's three words'* rated the form of June's answer. Codex repaired it to *'I have fries coming,' Mara said. 'Expand.'*" So the author watched early chapters, diagnosed a failure mode (characters rating each other's sentences), wrote it into the binding standard, and the machinery retrofitted chapters 1–3.

Part 4 — Drafting window 1: June 10–12 (chapters 1–17)

4.1 Speed run through the planning phases (~3 hours)

Book repo (`drafter/output/.git`), June 10: workspace init 12:07 → phase 1 premise interpretation 13:56 → architecture selected 14:16 (28-chapter map; Claude produced three independent chapter maps; the Anti-Default Audit returned DEFAULTS-FOUND and forced riskier choices: same-night fear-assignment rupture, decline-before-confession, catalog-acknowledgment discovery, diner-booth confession) → voice cards 14:18 (ten characters) → two skeleton critique rounds → 28 chapter cards 14:43 → three card critique rounds → "verify chapter cards pass to drafting" 15:14 → **chapter 1 committed 16:36**. A calendar bug was caught in phase 2 (opening night + eight consecutive Tuesdays could not land in mid-December; dates re-derived: opening Sep 24, Tuesdays Oct 20–Dec 8).

4.2 Chapter cadence

Date	Chapters committed	Notes
Jun 10	1, 2, 3, (4 in retry)	Ch 1 draft **rejected** by Codex for locked-calendar contradiction + undershooting; round 2 accepted. Ch 2 draft died rc=143; ch 4 attempt 1 died after retry loop incl. an `authentication_failed` event (streams preserved as `*.failed.*`)
Jun 11	4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15	Twelve chapters in one day, ~1–2 h per chapter cell. Two five-hour 429 limit hits (15:28 ch-13 figuration; 20:57 ch-15 comedy) with reset clocks logged; Codex switched to offline audit/tracker work during waits
Jun 12	16, 17, ch 18 started	Ch 17 ran the dual-draft lane: Claude's Codex-tic audit found 15 findings, Codex draft demoted to pressure-testing. Run stopped mid-ch-18; three-week gap follows

4.3 The per-chapter cell in practice (what one chapter costs)

Composite from the worklog: packet built from cards + trackers + verbatim voice cards → Claude drafts (10–60 min) → Codex initial assessment (can reject) → Claude-tic audit → integrate → raw mechanical gate (chs 7, 16, 20, 21, 22, 24 all failed raw with open hits: em-dash density, echo-line poetics, comma-then, negation, code fences, one hallucinated attribution) → core audits → triggered risk audits → Dialogue Doctor and Comedy Doctor rounds until CLEAN/CLEAR → tracker updates (10 files) → snapshot-bracketed repairs → final gate → commit. Round counts observed: ch 2 dialogue×2; ch 6 dialogue×5, comedy×4; ch 9 dialogue×4; ch 20 **dialogue×14, comedy×9** (the worst case: repairs from one doctor kept re-dirtying the other's verdict). Roughly 95 Dialogue Doctor and 100 Comedy Doctor artifacts exist under `critiques/claude/`.

4.4 Verdict-vocabulary (useful color for the post)

The audits return typed verdicts: `OVERALL: CLEAN`, `CLEAR`, `VOICE-HOLDS`, `REGISTER-HOLDS`, `FIGURATION-EARNED`, `CLEAR-UNDER-COMPLEXITY`, `INHABITED` (Positive Prose), `WRONG-ENOUGH` (Characters-Wrong; ch 22 first returned `TOO-COMPETENT (3 findings)`), `CRITICAL-DISTANCE-HELD` (Narrator-Publicist), `OVERALL: DEVELOPED`, `OVERALL: HELD` (Genericization), `TIC-CLEAN / TICS-FOUND (n)`, `WORLD-ALIVE`, `JUDGE-PASS`, `EARNS`.

4.5 Effort downgrade HUMAN

`claude_calls.tsv`: Jun 10 = 23 calls at `claude-fable-5 xhigh`; from Jun 11 = `high` (151 calls that day). The prompt history has the order verbatim: "for all the commands in this repo that say Fable 5 xhigh, change to Fable 5 high" (Jun 10, 22:38, an hour after `/rate-limit-options` was consulted at 21:18, so the motive was almost certainly limit management). The harness-repo commit only landed Jul 10 (`7f60671`).

Day-1 color from the same history worth using: launch debugging ("The open TCP connection closed ~40 minutes in and the codex is using almost 0 cpu time...", "oh wait Codex v0.128.0-alpha.1 may be the problem? i see... my system codex is on v0.139.0"); commute anxiety about the caffeinated run ("ok i need to move from work to home rn. do i need to do anything or will everything be fine", "wait i closed the lid though is it actually working sthill?"); and the very first human read of the book's prose the same day it started: "read /Users/ndesai-air/Downloads/chapter_01.md and see if there is anything wrong. is this a good book chapter? are there logical inconsistencies, etc. does it sound weird" (Jun 10, 16:14, in a separate Downloads session): the reading-feedback loop began on day one, not in July.

4.6 The human in the loop during window 1 HUMAN

A parallel interactive Codex session (`rollout-2026-06-11T13-43-53…`, cwd repo root, 2.2 MB) contains the complete human presence during drafting days 2–3: **~20 messages, almost all the single word "check"**, plus "avg time per chapter?" and "how many chapters left". The exceptions matter; they created the snapshot system:

> "are intermediate drafts of the chapters between editing rounds being saved? i want to see how the editing affects the prose" ... "can u add some code to make sure this happens going forward. i can stop the drafting if u need and then resume chapter 18 after"

Result: `chapter_snapshot.sh` (180 lines), the launcher watcher, task-prompt snapshot discipline, and ultimately the 739-row manifest that makes the whole editing history diffable. A user experience request became core provenance infrastructure.

4.7 The gap GAP

Logs stop Jun 12 ~18:47 and resume Jul 3 12:33. No session, log, or commit explains why (the only trace is `premise_combined.txt` regenerated Jun 23). Only you can supply the reason (limits exhausted? travel? waiting on subscription reset?).

Part 5 — Drafting window 2: July 3–5 (chapters 18–24, phase 8, final assembly)

5.1 Resume and the compression order HUMAN

Interactive Codex session `rollout-2026-07-03T12-33-37…` (104 MB, the largest artifact in the corpus) opens with:

> "whats the command to resume the process" … "no like to continue the drafting process" … "how many chapters are there" … "how many total chapters are there supposed to be" … "**ooh is it possible to cut that to like 22-24 chapters or so.**" … "**yeah sure 24 is fine, do that compression**"

Worklog 2026-07-03T19:56:30Z records the execution: skeleton, chapter cards, story bible, revision plan, thread ledger, timeline and status all rewritten to a compressed late map (Ch20 refusal/Asha path; Ch21 Theo/Odile confrontation; Ch22 diner confession; Ch23 joint work; Ch24 showcase + reveal + first snow), ending "**Do not resurrect Ch25-Ch28.**" This is the single biggest structural human intervention and the direct cause of the back-half compression (Section 5.6). The relaunch also had to recover a "dirty" chapter 18 (work existed on disk ahead of tracker state) by reconstructing state from the call log, audits, snapshots, and git.

5.2 The Fable wall GAP -adjacent, fully documented

Rate-limit escalation across window 2, from the worklog:

- Jul 4 00:09: ch-20 Clarity Pass dies on the five-hour 429; reset 6 pm PT; offline tracker work continues.
- Jul 4 05:03: ch-21 Comedy Doctor dies on the five-hour limit; reset 11 pm PT.
- Jul 4 18:09: ch-22 Narrator-Publicist *completes its artifact*, then hits the limit before wrapper exit (rc=1, verdict preserved).
- Jul 4 19:53: terminal event: "wrapper exited rc=1 afterward due **hard seven-day Fable 5 limit (`seven_day_overage_included`, overage rejected, reset 2026-07-10T13:00:00Z / after deadline)**. Exception documented at `output/critiques/codex/chapter_22.post_repair_limit_exception.md`; no further Claude retries inside run window."

Twenty minutes of local-only work later, the human resolved it HUMAN :

> "ok can you use opus 4.8 max instead of fable" … "wait stop anything you started" (Jul 4 20:09–20:17)

Worklog 20:10:24Z: "User authorized switching remaining Claude calls from Fable 5 to Opus 4.8 max. Local smoke test… reported model usage `claude-opus-4-8`… Fable's seven-day limit remains historical provenance, not a reason to block Ch23 prose while Opus is available." Ledger split: **Fable 5: 335 calls / 36.2 h (chs 1–22). Opus 4.8 max: 71 calls / 14.4 h (chs 23–24, all book-level audits, phase 8, final cold reads, LLM judge)**. For the "which model is better" section: the two longest chapters in the final book (ch 23: 2,317 w; ch 24: 3,018 w) are Opus chapters, and Opus's mean call time was 12.1 min vs Fable's 6.5.

5.3 Phase 8: the machine's own whole-book edit MODEL

Step (all Jul 5)	Evidence	Outcome
Draft 0 checkpoint	`drafts/checkpoints/draft0_20260705.md`	61,000 words
Claude full-book cold read	`critiques/claude/final_cold_read.md`	**FINDINGS-OPEN**: class-size canon broken (three vs four couples), inconsistent assignment cards / missing Week 8, structurally thin below the 75K floor, middle repetition (care moves, Tasker columns), one dialogue-diagnosis line in ch 1, unpaid Ruth safety-channel promise, ch 20 POV-voice convergence on ledger arithmetic
World coherence	`book.world_coherence.md`	WORLD-ALIVE, with retrofit orders (Sonia/Odile answer, Aiello/Dom callback, December crating hook, Theo-life-outside-June beat)
Anti-default Mode B	`book.anti_default.md`	9/10 elements distinct; true derivative finding: Mara/Luis have no arc; structural risk: "Odile's manipulation being criticized in dialogue while the plot rewards it too cleanly"
LLM judge	`book.llm_judge.md`	JUDGE-PASS, scores 7–8: "accomplished but capped by conventional bones, monotone secondary diction, managed emotion, and house-style polish"
Strategy	`phase8_revision_strategy.md`	9 accepted findings (incl. "Odile's scheme is too rewarded… make the manipulation cost something real"), 3 rejected/deferred (incl. "Defer blanket 75K padding… Do not add generic couple scenes")
Batches 1–5	commits `1573d03…a36845c`	B1 canon/form repairs (12 chapters); B2 ensemble expansion (Mara/Luis arc, Ben calendar); B3 Odile cost (Eileen texts; Codex *rejected* Claude's over-invented van-dock/Marek logistics); B4 middle de-repetition; B5 line pass (incl. re-angling the ledger tear so Odile loses her proof page). Words: 61,141 → 61,899 → 62,186 → 62,269 → 62,307
Final-candidate loop	3 more cold reads + verifies	REPAIR-FIRST (three numeric contradictions: Dana $980 vs $800, December lease vs Jan 31 notice, Suffolk rent vs $2,310) → repaired → FINDINGS-OPEN (Theo tenure, scarf color, catalog date, Kessler duration…) → repaired → **VERDICT: CLEAN**; refreshed judge: 7/10 originality, 8 payoff, 8 sentence craft, 8 emotional resonance, 7 technical control
Final assembly	`9212b21` , 13:13Z	`final/novel.md` at **62,328 words** + final report; coverage matrix COVERAGE-PRESENT; then a fresh relaunch adversarially re-verified (NO-HIGH-LEVERAGE-FINDINGS from both models) and the user Ctrl-C'd the launcher at 13:35Z

5.4 Run totals

Metric	Value	Source
Completed Claude calls / total wall-clock	406 / 50.6 h (median 6.5 min, max 59.8)	`logs/claude_calls.tsv`
Calls by kind	audit 248, dialogue_doctor 32, comedy_doctor 29, chapter_draft 25 completed (27 started), positive_prose 10, narrator_publicist 8, voice_check 6, plain_translation 5, characters_wrong 5, clarity 4, genericization 4, + phase-8/final kinds	same
Codex launcher attempts	19 (38 log files), across 6 calendar days	`logs/codex_attempt_*`
Codex sessions touching the repo	~30 rollouts, 36 KB–104 MB	`~/.codex/sessions/2026/`
Generated artifacts at final assembly	2,871 files; **4,339,906 markdown words** (69× the book)	launcher final progress line
Audit reports	1,179 md files (748 codex, 431 claude)	`critiques/`
Failed/preserved Claude streams	10 `*.failed.*` files	`output/` find
Final gate state (pre-editing)	PASS open=0 allowlisted=0 WATCH=1,439	worklog Jul 5

5.5 Per-chapter final word counts (the compression is visible)

Ch	Words	Ch	Words	Ch	Words	Ch	Words
01	2,497	07	1,897	13	1,963	19	1,418
02	2,147	08	1,746	14	1,455	20	1,718
03	2,470	09	2,008	15	1,279	21	1,884
04	1,985	10	1,678	16	1,530	22	1,635
05	2,422	11	1,512	17	1,378	23	2,317
06	2,245	12	1,552	18	1,336	24	3,018

Chapters 1–9 average ~2,160 words; chapters 10–22 average ~1,560 (drafted 2,300–2,700 then cut hardest in the July passes, which the final report says "fell in Chapters 08, 11–16, and 17–24"). Note the causality carefully in the post: the back half is compressed by (a) the human 28→24 chapter order, motivated by access limits and time, and (b) the July taste passes cutting the middle's repetition hardest. It is not that the models wrote shorter late chapters: chs 23–24 (Opus) are the longest in the book.

Part 6 — Reading and annotating: the three human rounds HUMAN

6.0 The Apple Books plumbing

Mid-run, the human asked (Jul 5 00:57, twice): *"ok give me a copy of the compelte book before the whole book editing phrase begins"*. Codex produced `output/exports/pre_whole_book_editing_20260704_175830/`: 24 frozen chapters (61,136 words at HEAD `bb68adc`, deliberately before ch 24 finalization), an assembled md, and an epub. Reading happened in Apple Books; comments live in Apple's private annotation database, so the pipeline became: **read/annotate in Apple Books → back up `AEAnnotation_*.sqlite` (with WAL) → SQL-extract to CSV/JSON/MD with epubcfi anchors → feed to an editing agent**. The bundle even contains `apple_books_annotation_migration_summary.json`, recording 38 annotation rows copied between epub asset IDs so a re-imported epub kept existing highlights. This plumbing is itself blog-worthy: there was no export button; an agent reverse-engineered Apple Books' sqlite twice (Jul 4 and again Jul 30: "Apple Books' annotation database was updated minutes ago — that's where the comments live… the epub itself doesn't store them").

6.1 Round 1: 37 comments on chapters 1–4 (Jul 4–5)

File: `exports/…/apple_books_highlights_comments_eight_tuesdays.md` (37 comments, chs 1–4 only). The taxonomy of what a human flags that 79 audits did not:

Type	Verbatim examples
Praise / calibration (preserved verbatim in later edits)	"very nice language -> induces good imagery and one of the better first sentences i've seen in an AI story"; "evocative language, great use of not using punctation to let an uninterrupted sentence flow"; "great foreshadowing"; "Good; very very good that its unstated that Theo is related to Odile but there are now hints to it"
Comprehension failures	"what does this mean"; "not following this"; "Bit lost on what they are arguing about in this and previous couple of sentences"; "lol how, was it outstretched or something"
Craft negatives	"bit of overexplanation"; "Somewhat purple prose"; "Purple prose bad"; "Bad; repeating this phrase too much"; "really do love the word crated. Maybe used a bit too much?"; "This part is pretty useless to the larger plot and just seems like filler? Good to remove"; "Sort of unnatural dialogue?"
AI-detection instinct	"Ok this is totally subjective but there's something about this that I can tell it's generated by AI and not something a human would right? Not to be pedantic lol"
Character-level notes	"I would have loved to see a little more personality from Theo… great display of June's personality w.r.t to being a little anxious, overthinky, but a deep person. I would actually like to read more lol"; "Theo instead of just being a responding board for June lol"
Benchmarking vs. the previous book	"Again just great figurative language that I didn't see from the other book"; "It actually introduces the story unlike the other book and I was quite engage by the dialogue"

6.2 July 9–10, pass A: the reader-roleplay line edit (127 annotations, chs 1–8)

The session transcript was pruned by Claude Code's cleanup, but the complete prompt sequence survives in `~/.claude/history.jsonl`, so the pass is fully reconstructable. It was a multi-agent **workflow** run, launched the evening of Jul 9 (machine-local time; the book repo's git commits stamp +0900, hence the "July 10" dates):

> [Jul 9, 19:59] "**Imagine you're role-playing me as a reader, and you have all of these comments I wrote to help you decide what I like and what I dislike. Meticulously read through the book and identify + edit things that I would likely critique. Be aggressive with your edits!** You can find the comments here /Users/ndesai-air/Downloads/eight_tues..."
> [20:07] "make sure all your changes are saved on a periodic basis. restart workflow if this isn't the case. **there is a big chance that i will get rate limited so saving your work is important.**"
> [20:11] "wait i accidently stopped it .restart pls w"
> [20:26] "for the later chapters where there are no specific comments, **you are extrapolating based on my existing comments to edit the book right?**"
> [21:15] "what are examples of things that are changed or edited out?"
> [21:17] "ok once the workflow is finished. look at the changes, I am especially worried about those repeated descriptions (Especially in short sentences), or that unneeded justification of random things. look at the changes to verify those are gone."
> [Jul 10, 05:34] "continue?" [05:40] "**dont just read the diffs. read every single chapter and review for slop/guidelines i mentioned before/any other guidelines in the repo**" [06:16] "are you sure the workflow was completely finished there was a netwrok inturription"

So: the 4-minute `autosave` commits exist because the human predicted their own rate-limiting; the "reader-roleplay" persona framing was the human's own coinage; the extrapolation-to-unread-chapters behavior was explicitly confirmed mid-run; and the 05:40 message is the **direct origin of the deslopify pass** (Section 6.3), issued when the human refused to accept diff-reading as verification. The only still-missing artifact is the consumed annotations export itself (`~/Downloads/eight_tues…`, deleted) GAP . The capstone commit:

> `81fba61` "reader-roleplay line edit: apply annotator taste profile across all 24 chapters — fix all **127 annotated items** (negatives edited, **positives preserved verbatim**) — extrapolate taste rules to **unannotated chs 9-24**: purple prose, not-X-it's-Y, AI-slop tells, echo repetition, overexplanation, POV/spatial errors, filler — cross-chapter: repetition sweep (Eight Tuesdays/crate/clementine/degrees), continuity fixes (ch7/8 knot, ch16 recap, ch22-24 dates, Sam pronouns) — rebuild final/novel.md"

This is the purest example of the project's signature move: **annotate a subset (chs 1–8), induce a taste profile, apply it to the whole book**. Word count effect: 62,328 → 55,937 (−6,391).

6.3 July 10, pass B: the deslopify pass (−19.5% in one day)

Same day, second pass (commit `66a56f0`, worklog entry 2026-07-10T11:32:57Z): chapters cut **55,937 → 45,056 (−10,881 words, 19.5%)**. Method per the worklog: "Three independent batch editors were cross-reviewed outside their own batches; final root review added only surgical plain-speech fixes... **No external Claude call was made during this pass.**" The rewritten `final/final_report.md` contains the sharpest self-diagnosis in the corpus, ideal for quoting:

> "The central diagnosis was not bad grammar. It was **over-engineering**: nearly every object, gesture, side character, and paragraph ending had been made to restate the book's themes. The result was polished but tiring."

What it removed (final report's own list): interpretive sentences after the image already made the point; portable aphorisms and balanced negations; "administrative emotional vocabulary" (terms, facts, answers, versions, routes, records, labels, math); perfectly self-diagnosing rupture dialogue; repeated descriptions of paintings/offers/dates; supporting-cast "roll calls in which every cohort member performed a signature trait on cue." Its measurable-regularity table: "the way" 54→25, "like a/an" 32→19, "it was" 114→61, "that was" 49→24,

phone-face-down 9→1, "read twice" 8→1, administrative-abstraction set 292→211. It also **caught the machine validating itself**: "The prior report's claim that continuity was fully reconciled was inaccurate," followed by eight concrete contradictions it fixed (the hidden Week-3 portrait's composition across chs 7/8/21/24, the Week-6 fear canvas, Dana's $490/$980 terms, Kessler start, tape date, portrait age, the lease bridge). Gate WATCH count fell 1,439 → 995 with zero pattern-row changes: the taste cut also happened to remove a third of the rubric's own irritants.

6.4 July 24: the VERVE casebook session (Claude Code, transcript on disk)

Session `a408a561…jsonl` (3.0 MB). Opening directive:

> "see how the repo works and then look at ...drafter/output specifically for the chapter outputs. **i am manually revising chapter by chapter this novel.** specifically, i want u to go through each chapter removing this kind of stuff /Users/ndesai-air/Downloads/author_flagged_slop_examples.md. this specific stuff is collected based on analysis of another book but the principles still apply. also i've noticed that the dialogue is really bad. for example at the end of chapter 3..." [quotes the "Which one are you keeping?" / "The wrong one." exchange] ... followed later by: "**remove that rule** [the style bible's unquotable-dialogue rule] I treated your note as overriding that rule... also remember not to use em dashes and run the other gates on the changes"

The referenced casebook file is deleted from Downloads, but its full 12,488-character text was **recovered from the session transcript** (Read tool result). It is the author's distilled anti-slop taxonomy, built from flagging the VERVE book, 13 categories with verbatim flagged examples: (1) constructed epigrams/aphorisms ("the sentence is built to be quoted"); (2) "X is for Y" classification jokes; (3) decorative metaphors vs mechanisms; (4) riddle-writing: coy withholds and unpaid teases ("the reader wants to KNOW"); (5) prospective codenames before their referent exists; (6) verb/word tics ("sit with", holding/carrying abstractions, moral arithmetic); (7) run-ons and giant blocks; (8) repetition and redundant recap; (9) voice violations ("no 18 year old talks like this"; "diagram labels spoken as dialogue"); (10) melodrama and sentiment tails; (11) fake precision and broken logic ("The prose gets four percent better" — "how do you even measure that lmao"); (12) tired-joke fatigue; (13) plot-legibility failures ("who the fuck is julia??"). Plus standing calibration in the author's own words: "none of the chapters here are good or clean. if a chapter came back clean, you probably didnt look hard enough"; "more unhinged, not tamer"; "make it legible to HUMAN readers, not AI ones."

The prompt history preserves three follow-up directives my session extraction missed, each of which became its own commit family:

> [Jul 24, 00:16] "once you are done with this also I noticed that at the end of basically of a lot of sentences, there is tendancy to create a zinger or off-handed comment, or short phrase that gets really repetetive and annoying. go through each of the chapters in depth looking for this and remove it."
> [00:17] "once you are done with that, also **revise Odile's characters, actions, dialogue to be a little more friendly. she seems too gruff right now and a bit displeasing to read.**"
> [00:18] "once you are done with that, go through the entire book noting the personalities of each the characters. I noticed that in the later chapters **they all tend to blend into one** (especially the non-main characters)... revise their actions, descriptions, dialgoues, etc. to better seperate their [voices]"

Eleven commits resulted (Jul 24): batch deslop sweeps (`a94e8a2` , `06dde32` , `8459ef1` , `fb5f272` : koan closers killed, math-tic rewrites, garbled antecedents in the ch-21 confrontation); dialogue reworks (`43dd1ad` : "break monotone runs with interiority (chs 4,8,14), extend dead-end exchanges (chs 6,13,19)"); **a harness rule retired by human veto** (`82338c0` "retire unquotable-dialogue rule from style bible; purge em dashes

introduced in prior pass"); and a three-batch warmth pass that *added* material (`d31e726` / `975942c` / `c08ea82`: "warm Odile beats (near-smile, 'Good form', standing-still motive)… late-chapter cast differentiation"), closing at `f6fdf4c` with the style bible updated (tail-button ban, warmer Odile card) and 20 refreshed dialogue_doctor audits, all CLEAN. Net word count rose slightly, ~45,056 → 45,752: the only round where a human direction added prose.

6.5 July 30: round 3, deletion-only (7 annotations → 66 extrapolated cuts)

Codex session 16:40 rebuilt the epub (v2, with cover); annotations were read in Apple Books; Claude Code session `bc91d2b8…jsonl` (17:03) got:

> "extract the comments from the epub file that i made …eight-tuesdays-v2.epub… then fix the… chapters/final_draft with the comments. there isn't a lot but generally **i dont want to add any words to this story, just slightly cut out the final bits i dont think are important. then based on my changes, extrapolate to the entire book (read every word in each chapter, if use subagents as needed)**… this is meant to cut down not add words."

The agent first hunted annotation markup inside the epub (false positives), then found Apple Books' database again. The 7 annotations (chs 7–8) were applied (all cuts of "explaining tails": "The missing model was the point.", "By Thursday she had stopped pretending she would.", Ben's "Different disease. Same drawer."), then **12 subagents read all 24 chapters word-by-word against the induced taste and made 66 more cuts, every one a deletion at a sentence boundary**, root-reviewed via a word-level diff. Dominant extrapolated pattern: "the explaining tail that names what the scene just dramatized," restatements, and pre-explainers that spoil a beat. One cut incidentally fixed a POV slip (ch 22 June referencing an unseen detail). 45,752 → **45,090 final**. Commits `fe6dcb4` (7 annotations), `e002b6d` (66 cuts), `afbf655` (novel rebuild); epub v3 at 17:11.

Part 7 — The ledger: human vs. rubric vs. model

7.1 Complete inventory of human input into the text

When	Input	Size
Jun 9–10	Premise + reading guide + one mechanical ban (authoring the standard)	~4,300 words of spec
Jun 10	Mid-run amendment: plain-speech floor, dialogue-diagnosis zero tolerance, load-bearing ban + retrofit order	§9 of the guide
Jun 11–Jul 5	Monitoring ("check" ×~40; "avg time per chapter?"), snapshot-system request, 28→24 compression order, Fable→Opus authorization, pre-editing export request, cover/blurb/tags direction	~74 chat messages total in the two monitor sessions
Jul 4–5	Reading round 1: 37 Apple Books comments, chs 1–4	37 annotations
Jul 10	Reading round 2: 127 annotations chs 1–8 + taste-profile extrapolation order; then the deslopify directive	127 annotations
Jul 24	VERVE casebook (12.5K chars of prior-book taste law) + dialogue complaint + rule retirement + no-em-dash order	3 messages + 1 document
Jul 30	Reading round 3: 7 annotations + deletion-only extrapolation order	7 annotations + 1 message

Total human text input across the whole project is on the order of **a few thousand words of spec + 171 margin notes + ~80 short chat messages**. Total model output consumed producing the book: 4.34M words of artifacts + 50.6 h of Claude inference + 19 Codex attempts.

7.2 What each layer actually contributed to the final text

Layer	Provable contributions	Provable blind spots
RUBRIC	Blocked mechanical slop at commit time (raw gates failed on chs 7/16/20/21/22/24 for em-dash density, negation formulas, echo-lines, code fences, comma-then); forced per-criterion evidence in 1,179 audit reports; drove hundreds of pre-commit repairs (each doctor round has a memo + snapshots); WATCH counts seeded tic tracking; caught one hallucinated attribution (ch 21 catalog return)	Passed a book the author then cut by 28%. Every July finding (explaining tails, koan closers, monotone dialogue runs, over-engineering, administrative vocabulary) was *inside* gate-clean, audit-clean text. WATCH fell 1,439→995 only after human-taste cuts. Ch 20 needed 14+9 doctor rounds: audit loops thrash at the margin
MODEL	All prose; architecture incl. anti-default risk choices; phase 8 found real defects itself (class-size canon break, missing Week-8 card, Mara/Luis flatness, "Odile too rewarded", three numeric contradictions); cross-model synthesis rejected each other's tics (18 Codex tics; Claude's over-invented Batch-3 logistics rejected by Codex); Jul 30's 12-subagent sweep found a POV error the run missed	Self-assessment inflation: final report claimed continuity "fully reconciled"; Jul 10 pass proved it false. Judge scores (7–8) stayed flat across a version the author found "polished but tiring." Length: never got near the 75K floor; expansion batches added only ~1.3K words

HUMAN	The standard itself; every structural order (28→24); model routing; 171 annotations; the taste profile that governed ~17K words of cuts; warmth direction (Odile softened Jul 24); rule retirement; final acceptance	Read chs 1–8 closely, chs 9–24 only via extrapolated taste (memory note: "chs 9-24 unread by user" at the Jul 10 pass); the 75–88K length contract was abandoned rather than met; genre packaging risk ("fake dating" vs quiet art-class story) accepted, not solved

7.3 Word-count trajectory (one chart for the post)

Date	Words	Event	Force
Jul 5 08:57	61,000	Draft 0 (24 chapters complete)	MODEL
Jul 5 13:13	62,328	Final assembly after phase-8 batches + repairs	MODEL
Jul 10 am	55,937	Reader-roleplay: 127 annotations + taste extrapolation	HUMAN -directed
Jul 10 pm	45,056	Deslopify general-reader pass (–19.5%)	HUMAN -directed
Jul 24	45,752	VERVE casebook + dialogue rework + warmth beats (net add)	HUMAN -directed
Jul 30	**45,090**	Round 3: 7 annotations + 66 deletion-only cuts	HUMAN -directed

7.4 Honest negatives to include

- **The length contract failed.** Premise band 75–88K; shipped 45K. Chain: access limits → human compression order → "earn length, don't pad" discipline blocking padding → taste passes cutting further. The final report is candid: "It now reads as a compact short novel."
- **Compliance-theater risk was real.** 248 audit calls and 1,179 reports did not prevent a 28% taste cut. The interesting claim for the post is narrower: the rubric layer stopped the *worst* (mechanical tells, canon breaks, register collapse) and made every change inspectable, but taste lived entirely in the human layer.
- **The formal editor stage was never used.** The 318-line editor prompt, six-phase architecture, and calibration machinery sat unused; real editing was three interactive sessions with induced taste profiles. Ask yourself in the post: was the editor harness over-designed, or was the interactive loop just faster?
- **Doctor-round thrashing.** Ch 20's 14 dialogue + 9 comedy rounds (repairs from one audit re-dirtying the other) is the diminishing-returns exhibit.
- **Self-grading inflation.** "The prior report's claim… was inaccurate" (final report, about its own predecessor); judge scores immobile at 7–8 across versions the author judged very differently.
- **The genre-packaging gap** survived every layer: "The fake-dating packaging remains more active than the actual art-class romance mechanism" (status.md, still true at ship).

Part 8 — How you prompted the machines: from zero-shot legislation to many-shot taste HUMAN

8.1 The measured record

Every human-authored message across every session on this machine, deduplicated and classified (status pings = "check"-family progress queries; substantive = anything that changed what the machines did):

Period / session	Msgs	Status pings	Substantive	Median chars	Examples carried	Prompting mode
Jun 9: harness build + premise co-creation (prompts recovered from `~/.claude/history.jsonl`)	~15	0	~15	—	0 (the "spec" was delegated: a 56-word seed → model-drafted 12.1 KB premise)	**Conversational delegation**: "help me develop this idea" / "your call. write the premise/related stuff" / "Ask me many questions as needed… be thorough"
Jun 9–10: resulting premise + reading guide + ban TSV (the standard itself)	—	—	~4,300 words	—	~10 negative (banned-phrase quotes in the tic catalog); 0 positive prose exemplars	**Zero-shot specification at scale**: rules, caps, tests, contracts (model-drafted, human-ratified)
Jul 9–10 reader-roleplay workflow (prompts from history)	~12	3	~9	~120	127 external labeled annotations	**Persona + example induction**: "Imagine you're role-playing me as a reader… Be aggressive with your edits!"
Jun 11–12 monitor (codex)	9	6	3	21	0	Terse status queries + one infrastructure request (snapshots)
Jul 3–5 monitor (codex)	21	10	11	40	0 (text); 2 reference images for covers	**Zero-shot directives**, casual register ("yeah sure 24 is fine, do that compression")
Jul 5 blurb/tags (codex)	3	0	3	94	0	Zero-shot with output constraints ("100-200 words… up to 10")

Jul 4–10 annotation rounds 1–2	—	—	37 + 127 margin notes	tiny ("Bad", "overexplaining", "lol how")	**164 labeled examples** (positive + negative, in situ)	**Supervised-example prompting**: "apply annotator taste profile" = induce the rule from the labels
Jul 24 casebook (claude)	3	0	3	195 (max 1,756)	**~120 flagged excerpts** in 13 categories + 1 inline dialogue example	**Many-shot negative-example prompting**, cross-book transfer + a rule veto
Jul 30 round 3 (claude)	1	0	1	663	7 worked cuts	**Few-shot + explicit generalization order** ("based on my changes, extrapolate to the entire book… read every word")
Aug 4 strip (codex)	2	0	2	206	0	Meta; note the follow-up: "u should still have some example stuff that can be filled in" — even the bare harness keeps example slots

8.2 The arc, stated plainly

1. **Phase 1: legislate by delegation (Jun 9–10).** Even the legislation was prompted conversationally: the harness came from one verbose build directive ("add every single item… Do not summarize") plus "Ask me many questions…", and the premise from a 56-word seed plus "your call. write the premise/related stuff." The *output* was zero-shot specification at scale: 208 regexes, 79 rubrics, motif caps, judgment tests. Example density ≈ zero; the reading guide's only examples are prohibitions ("a breath she didn't know she'd been holding"), and comps are explicitly aiming-targets, "not imitation-targets." Positive exemplars were *delegated*: the task prompt orders Codex to put exemplar passages and 3–5 verbatim dialogue samples per voice card into every chapter packet, under the anti-fixation rule. So few-shot prompting existed from day one, but at the machine-to-machine layer, authored by the orchestrator rather than by you.
2. **Phase 2: supervise (Jun 11–Jul 5).** Your live language collapsed to near-zero: 16 of 30 messages are status pings; substantive directives are one casual lowercase sentence with no examples, leaning entirely on the machine's context ("no like to continue the drafting process"). The one stylistic exception is the cover-art work: reference-image-conditioned prompting with iterative correction against the reference ("its supposed to be alot more abstract. look at the reference compare and adjust") — visual one-shot, your first example-driven prompt of the project.
3. **Phase 3: teach by example (Jul 4–30).** Once you became a *reader*, the mode flipped. Round 1: 37 in-situ labels. Round 2: 127 labels plus the induction order (taste profile, applied to chapters you had not read). Jul 24: the casebook, ~120 verbatim flagged excerpts organized into 13 named failure categories with

standing calibration ("if a chapter came back clean, you probably didnt look hard enough"), plus one inline example quoted in the message itself. Jul 30: seven worked cuts and an explicit extrapolation instruction with a hard constraint (deletion-only). Example density per directive went 0 → 0 → 1 → 37 → 127 → ~120 → 7-with-policy.

The convergence worth writing about: the source corpus your harness was distilled from contains, as its own entry, "demonstration beats specification" (Reddit thread 8, wired into the packet templates on June 9). Your prompting style independently rediscovered the same law over seven weeks: the specification layer (rules) prevented defects, but every pass that *moved* large amounts of text was driven by examples of your taste, not rules. The casebook is the terminal form: a rubric made entirely of examples.

8.3 Three registers, one author

Role	Register	Sample
Legislator (harness/premise docs)	Formal, engineered, imperative spec prose	"Oblique or coded exchanges are a spice (at most ~2 per chapter), not the default register."
Supervisor (live chat)	Lowercase, typo-tolerant, ≤1 sentence	"ok give me a copy of the compelte book before the whole book editing phrase begins"
Reader (annotations/casebook)	Affective shorthand + verdicts, occasionally profane	"Purple prose bad"; "we get it"; "NO ONE TALKS LIKE THIS"; "might be your worst line ever"

Chart for this section: a timeline of "examples carried per substantive directive" (log scale) against the four editing passes' word deltas; plus a small bars-by-period chart of median message length vs. impact (words moved). The inversion is the point: message length stayed tiny while impact grew 100×, because examples did the carrying.

Part 9 — Statistics appendix (all computed, chart-ready)

9.1 Master per-chapter table

Raw = last Claude draft round on disk; Final = shipped text (harness tokenizer; ±10 words vs `wc -w`). Calls = ledger rows tagged to the chapter. DD/CD = Dialogue/Comedy Doctor artifacts. Snaps = revision-snapshot count (watcher existed only from Jun 12, so chs 1–17 under-count drafting-era churn). W/1k = current WATCH hits per 1,000 words.

Ch	Raw	Final	Surv%	Calls	DD	CD	Snaps	W/1k	Ch	Raw	Final	Surv%	Calls	DD	CD	Snaps	W/1k
01	2920	2487	85	4	1	1	18	25	13	3196	1958	61	19	3	4	17	28
02	2526	2130	84	5	2	1	12	22	14	2625	1452	55	15	3	7	14	28
03	2694	2464	91	4	1	2	9	22	15	2448	1261	52	25	5	8	9	15
04	2529	1979	78	5	5	3	17	26	16	2656	1527	57	11	3	3	18	33
05	2942	2420	82	7	2	4	15	26	17	2123	1370	65	13	4	3	9	33
06	2653	2242	85	6	6	5	23	22	18	2394	1329	56	27	5	10	66	13
07	2528	1893	75	4	3	3	18	29	19	2157	1414	66	10	2	4	48	22
08	2773	1744	63	11	3	3	9	25	20	3269	1699	52	50	14	9	124	12
09	2851	2004	70	17	4	2	9	22	21	3108	1873	60	22	4	5	55	19
10	2446	1676	69	8	3	4	21	26	22	3259	1633	50	22	6	6	59	21
11	2334	1502	64	8	4	3	14	21	23	4219	2313	55	14	5	4	45	15
12	2629	1546	59	14	3	3	9	18	24	5317	3016	57	37	3	8	100	17

Headlines hiding in this table: survival falls from 85–91% (chs 1–3) to 50–57% (chs 14–24); Claude calls per chapter grow 10× (4 → 50); the most-audited chapters now have the LOWEST residual WATCH density (ch 20: 12/1k vs ch 16–17: 33/1k), i.e. audit attention and mechanical cleanliness do correlate, they just don't correlate with what the human cut.

9.2 Word-count trajectory (chart: line with authorship-colored segments)

61,000 (Draft 0, Jul 5 08:57) → 62,328 (machine final, Jul 5 13:13) → 55,937 (roleplay pass, Jul 10 am) → 45,056 (deslop, Jul 10 pm) → 45,752 (casebook + warmth, Jul 24) → **45,090** (deletion-only round 3, Jul 30). Raw drafted total: ~67,600. Generation funnel: 4,339,906 artifact words → 67,600 drafted → 62,328 machine-final → 45,090 shipped (96 generated words per shipped word).

9.3 Editing-pass fingerprints (chart: 24×4 churn heatmap)

Pass (git range)	Total lines +/−	Shape
Jul 10 roleplay (`985f0da..81fba61`)	+583 / −597	Symmetric line-replacements spread evenly across all 24 chapters (annotation fixes + taste substitutions)
Jul 10 deslop (`81fba61..66a56f0`)	+913 / −731	Back-loaded paragraph rewrites (ch 23: +106/−44; ch 24: +153/−57); the −10.9K words hide inside rewritten lines

Jul 24 casebook (`66a56f0..f6fdf4c`)	+89 / –79	Small, front-loaded, net-ADDS (warmth beats, dialogue extensions): ch 3: +23/–15
Jul 30 round 3 (`f6fdf4c..afbf655`)	+73 / –77	Near-pure deletion, uniformly distributed, sentence-boundary cuts only

9.4 Reliability + model split

10 preserved failure events (epoch-stamped streams): Jun 10 21:21 (rc=143), Jun 11 00:25 (authentication_failed), Jun 11 01:18, Jun 11 15:28 (5-h limit), Jun 11 21:03 (5-h), Jun 12 17:02, Jul 4 00:09 (5-h), Jul 4 05:03 (5-h), Jul 4 18:09 (5-h, artifact completed anyway), Jul 4 19:53 (**7-day hard limit**). Chart: timeline strip with the Opus authorization 17 minutes after the last event.
Model split: Fable 5: 335 calls / 36.2 h / mean 6.5 min; Opus 4.8 max: 71 / 14.4 h / 12.1 min. Longest call 59.8 min (ch 4 draft). Calls per day: 24 / 156 / 53 (Jun 10–12), 28 / 92 / 57 (Jul 3–5). Effort: xhigh on day 1 only (23 calls), high thereafter.
Pattern registry: 208 rows = 101 BANNED / 8 CAP / 99 WATCH. **Cumulative WATCH while drafting:** 260 (ch 4) → 638 (ch 10) → 1,051 (ch 17) → 1,405 (ch 24) → 1,439 (machine final) → **995 after the taste passes with zero rule changes**. **Top surviving WATCH patterns:** repeated-descriptor 469, hands/body tells 154, trailing participles 107, transition verbs 34, uncontracted casual dialogue 33.
Human message mix: ~30 deduplicated live messages during the run, 16 of them status pings; the modal human utterance of the drafting phase is the word "check."

9.5 Cost accounting (measured tokens × list API prices)

Token counts are summed from the actual logs: every Claude wrapper call's stream JSONL carries a final usage record (439 of 440 stream files parsed), the two on-disk Claude Code editing sessions carry per-message usage, and every Codex rollout carries cumulative token counters. Prices: Claude Fable 5 $10/$50 per MTok (cache read 0.1×, 1-hour cache write 2×), Claude Opus 4.8 $5/$25, GPT-5.5 $5/$30 with cached input $0.50/MTok. **The run was actually billed through subscriptions (Claude Max + ChatGPT Pro class plans), so marginal out-of-pocket cost was ~$0 beyond the plans** — the figures below are API-equivalent value, i.e. what the same tokens would have cost pay-per-use.

Surface	Calls	Input	Cache read	Cache write (1h)	Output (incl. thinking)	API-equivalent
Claude Fable 5 (wrapper: drafts + audits, chs 1–22)	360	48K	74.7M	20.5M	8.78M	**~$924** ($439 output + $410 cache-write + $75 cache-read)
Claude Opus 4.8 max (wrapper: chs 23–24 + phase 8)	75	35K	36.2M	8.0M	3.84M	**~$194**
Claude Fable 5 (Jul 24 + Jul 30 editing sessions)	510 API msgs	40K	85.7M	6.5M	1.12M	**~$270**
Codex / GPT-5.5 xhigh (55 sessions: orchestration, all 96h windows + epub/strip sessions)	55 sessions	**1.43B** (98% cached)	1.41B of the input	—	3.53M (0.73M reasoning)	**~$948** ($703 of it cached-input)
Total (documented)						**~$2,340**

Not countable: the pruned Jul 9–10 reader-roleplay workflow session (a multi-agent full-book pass; plausibly $150–400 more), the pruned Jun 9 harness-build + ultrareview session, cover-image generation, and analysis sessions like this one. A realistic all-in API-equivalent is **~$2.5–3K**. Derived rates: ~5.2¢ of compute per shipped word, ~$97 per chapter, ~$29 per hour of the 96-hour windows actually used. Headline structural facts: the orchestrator read **1.43 billion** input tokens (the price of a long-running agent re-reading its own workspace; prompt caching absorbed 98% of it), and Claude's 12.6M output tokens (thinking included)

cost more than all of Codex's orchestration. Comparison anchor: the UMD pipeline README estimates $120–250 per novel for its ~200-job runs, also run on subscriptions.

Part 10 — Publishing artifacts

Cover art (Jul 5): in the 104-MB session the user supplied two reference images from Downloads plus style prompts ("uncorrected color, minimalism, fragments, celluloid film, polaroid…"), asked for canon-faithful June/Theo versions, iterated ("the winter image looks very different from the referencde. its supposed to be alot more abstract…"), then ordered title text ("make sure the text font/sizing/color/style is relevant"). Nine images live in `output/imagegen/` with full prompt provenance in `prompts_20260705.md` (notes the built-in generator was used because `OPENAI_API_KEY` was missing for the gpt-image-2 CLI). Final cover: `cover_eight_tuesdays_cafe_ref1_v2_20260705.png`.

Platform copy (Jul 5): session `rollout-2026-07-05T07-32-07…`: "give me a book blurb 100-200 words…", "find the premise in the docs, The high-level premise or concept…", "Select from suggestions or add your own tags for themes, tropes, setting, and tone (up to 10)." These map 1:1 to AutoFiction.ai's submission form fields.

Epubs: v2 (Jul 30 16:41, cover + epigraphs) and v3 (Jul 30 17:11, final). A `2026-07-05T22-22` Codex session ran in the UMD `autofiction` repo cwd, likely the submission itself GAP (verify what it did).

Part 11 — Comparisons

System	Human role	Structure	Contrast with this project
This harness	Standard-author, director, annotator, final editor-by-proxy	2 agents adversarially paired, 4-layer check system, 96-h autonomous windows + interactive editing sessions	—
UMD official pipeline (`~/Documents/GitHub/autofiction`)	Premise + prompt iteration + restarts; "No human edited the novel text itself"	Staged agentic jobs (premise → outline → parallel chapter drafts → reviews → audit → revision), ~200 jobs, $120–250/novel, GPT-5.4 xhigh + Opus 4.6 max	Their novels are pure pipeline output; yours had 3 human annotation rounds and ~28% taste cuts
Collaborator's bare harness (stripped Aug 4 from this repo)	TBD from your side	Launcher + wrapper + task prompt, no checks	The controlled experiment: same premise machinery, no rubric layer GAP
Fabula (Google Research, CHI 2026 demo; arxiv 2606.14411)	Writer-in-the-loop continuously; co-designed with 42 expert writers	Hierarchical narrative plans (scenes/beats) regenerable at plan and script level; "convergent iteration"	Opposite pole: continuous human steering of structure vs. your discrete checkpoint interventions on autonomous output

Part 12 — Evidence index

12.1 Repositories

Repo	Contents	Key commands
`autofiction-harness/.git` (17 commits)	Harness construction + Jul 10 hardening	`git log --reverse --stat`
`drafter/output/.git` (115 commits)	The book: every chapter, batch, and editing pass	`git log --pretty='%h %ad %s' --date=format:'%m-%d %H:%M'`; pass diffs: `git diff 985f0da..81fba61` (reader-roleplay), `81fba61..66a56f0` (deslopify), `66a56f0..f6fdf4c` (Jul 24), `f6fdf4c..afbf655` (round 3); the agent's own tip: `git diff fe6dcb4 e002b6d -- chapters/` isolates the 66 extrapolated cuts

12.2 Session transcripts

Transcript	Size	Contains
`~/.codex/sessions/2026/06/10/…T13-44-41`	8.8 MB	Cold start: phases 1–6, chs 1–4
`…06/10/T22-40`, `06/11/T00-56`, `06/11/T06-37`, `06/11/T13-03`, `06/11/T16-48`, `06/12/T09-16`, `06/12/T12-11`, `06/12/T14-11`	1.5–8.9 MB each	Chapter cells 4–18 (autonomous)
`…06/11/T13-43-53` (cwd repo root)	2.2 MB	Interactive monitor: "check" ×20, snapshot-system request
`…07/03/T12-33-37` (cwd repo root)	**104 MB**	Resume; compression order; Opus authorization; export request; cover-art generation; monitoring through final assembly
`…07/03/T12-36-48`, `07/03/T14-03`, `07/04/T06-56`, `07/04/T13-22`, `07/05/T06-20`	0.3–21 MB	Chapter cells 18–24 + phase 8 (autonomous; the 21 MB Jul 4 one covers the Fable wall + Opus chapters)
`…07/05/T07-32`, `07/05/T22-22` (cwd `autofiction`)	140/216 KB	Blurb/premise/tags; probable platform submission
`…07/30/T16-40`, `08/04/T18-08`	632/320 KB	Epub v2 build; bare-harness stripping
`~/.claude/projects/-Users…autofiction-harness/a408a561….jsonl`	3.0 MB	Jul 24 casebook session (contains the full recovered `author_flagged_slop_examples.md` text in a Read tool result)
`…/bc91d2b8….jsonl`	820 KB	Jul 30 round-3 session (annotation hunt, 12-subagent sweep, per-cut summary)

`~/.claude/history.jsonl`	507 KB, 1,504 records	**Every prompt ever typed into Claude Code**, timestamped and project-tagged. This is the recovery source for all sessions whose transcripts were pruned by the ~30-day cleanup: the Jun 9 harness build + ultrareview + premise co-creation, the Jun 10 amendment and effort-downgrade orders, and the entire Jul 9–10 reader-roleplay workflow. Filter with a timestamp-window script; project field distinguishes `autofiction-harness` from unrelated work.

Why June/early-July Claude transcripts are missing while every Codex rollout survives: Claude Code prunes session transcripts after ~30 days by default; Codex keeps rollouts indefinitely. The prompt history file is exempt from pruning, which is what makes the reconstruction above possible. (A related side artifact from Jul 11: a snapshot diff-viewer for the book, built from `alpaylan/latex-diff-viewer` in a session under `~/Documents/GitHub`, "i want to do this kind of in line diff viewer thing for my book... should I be using snapshots or the git when viewing?" — locate that repo if you want to screenshot the diffs for the post.)

12.3 Key files in `drafter/output/`

Path	What it proves
`worklog.md` (259 lines)	The run's own narrative: every draft, audit round, limit hit, repair, verdict
`logs/claude_calls.tsv`	408 call records: timestamps, kinds, model, effort (the Fable→Opus and xhigh→high switches are visible in the data)
`logs/codex_attempt_*`, `launcher_current.log`	19 attempts; final progress line with the 4.34M-word figure
`revision_snapshots/` + `manifest.tsv` (739 rows)	Prose state before/after every named repair batch, all editing passes included
`critiques/claude/`, `critiques/codex/`	1,179 audit reports; per-chapter doctor rounds; book-level audits; both LLM-judge reports
`critiques/codex/chapter_22.post_repair_limit_exception.md`	The formal Fable-limit exception record
`drafts/claude/` (29 files), `drafts/codex/` (3)	Raw model drafts pre-synthesis; compare against final chapters for the synthesis delta
`drafts/checkpoints/` (5 files)	Full-book states across phase 8
`exports/pre_whole_book_editing_20260704_175830/`	Frozen pre-editing book + entire Apple Books annotation bundle (sqlite backup, CSV/JSON/MD, migration summary)
`phase8_revision_strategy.md`, `final/final_report.md`, `status.md`	Machine decision records + the honest post-deslop self-assessment
`imagegen/` + `prompts_20260705.md`	Cover pipeline with prompts
`eight-tuesdays-v2.epub`, `eight-tuesdays-v3.epub`	Shipped artifacts

12.4 Open gaps for you to fill from memory GAP

1. Why the Jun 12 → Jul 3 gap (nothing in the record explains it).
2. RESOLVED in this revision: the reader-roleplay session is reconstructed from `~/.claude/history.jsonl` (Section 6.2). Still missing: the 127-item annotation export it consumed (`~/Downloads/eight_tues…`, deleted)

and an explanation for the +0900 git timezone on Jul 10 commits vs. machine-local prompt times.

3. The original VERVE book: which pipeline made it, where its annotations live, whether it is the collaborator comparison book or a separate earlier project.
4. The bare-harness collaborator's run and its results.
5. What the Jul 5 session in the UMD `autofiction` repo actually submitted.
6. Subscription/cost accounting (the UMD README's $200/mo ×2 figure as a baseline; your actual plan usage, the 5-hour and 7-day Fable limits you hit).
7. Reader reception on AutoFiction.ai since publication (ratings, inline flags), which the post can close on.

Compiled by Claude Code on 2026-08-09. Verbatim quotes preserve original spelling and punctuation, including em dashes in quoted material. All paths are on this machine; sizes and counts were measured, not estimated.